

## Appendix B: impute values for censored median\_house\_values

### Overview

Appendix B applies the imputation method set out in Appendix A to the records with censored median house values.

\* \* \* \* \*

### Section 1: get a prediction for the mean of the actual, unobserved values

In [ ]: *# Load some of the packages we will use.*

```
require(repr)    # allows us to resize the plots
require(stringr)
require(ggplot2)
require(car)     # needed for diagnostic tools
require(arm)
```

In [2]: `options(digits = 5, show.signif.stars = F,  
mc.cores=parallel::detectCores())`

In [3]: *# This dataset contains imputed values for housing\_median\_age.  
# The imputation was done in Appendix A.*

```
dat <- read.csv("/home/greg/Documents/stat/Geron_ML/datasets/housing/housing_cleaned_v03.csv",
               header=TRUE, row.names=1,
               colClasses= c("character", rep("numeric", 9), "character",
                               rep("numeric", 5)))
dim(dat)
```

20603 · 15

In [4]: *# Check that we have imputed values for housing\_median\_age.  
# Prior to the imputation done in Appendix A, the age values  
# were capped at 52.*

```
summary(dat$housing_median_age)
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.0	18.0	29.0	29.0	37.0	74.9

In [5]: *# Plot of median\_house\_value vs. median\_income.  
# 4.8% of the data is censored at 500K.*

```
options(repr.plot.width= 8, repr.plot.height= 6)

p <- ggplot(dat, aes(median_income, median_house_value)) +
  geom_point(alpha= 0.1) + xlab("median_income") + ylab("median_house_value") +
  ylim(0, 600000) +
  ggtitle("median_house_value vs. median_income,  
showing censored values at 500K") +
  theme(axis.text= element_text(size = 12)) +
  theme(axis.title= element_text(size= 14)) +
  theme(title= element_text(size= 16))
```

p



```
In [6]: # There are 990 records, or districts, with a
# censored median_house_value.

nrow(dat[which(dat$median_house_value >= 500000),])

990
```

## Get record counts for 15K interval bins of median\_house\_value

In order to mimic the age-level counts from Appendix A, we need to discretize median\_house\_value. For presentation purposes, I have chosen 15K rather than 10K for the interval size in order to reduce the variability in the counts.

For modeling purposes I make use of 10K bins because I want a prediction at 500K and 500K is not divisible by 15K. This lack of even division could potentially have a negative effect on the prediction at 500K.

```
In [7]: summary(dat$median_house_val)

   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
14999  119600  179800  206908  264950  500001
```

```
In [8]: # Let 15K be the lowest median_house_value in our dataset.

dat[which(dat$median_house_value < 15000), c("median_house_value")] <- 15000
```

```
In [10]: # Function for obtaining the number of records in each 15K
# interval.

get_rcd_counts <- function(med_houseVal, varRange,
                           span=15000, startpt=15000) {

  cell_floors <- seq(from=startpt, to=990000, by=span)
  names(cell_floors) <- paste(as.character(cell_floors/1000), "K", sep="")

  cell_floors_tmp <- cell_floors[(as.numeric(cell_floors) >= varRange[1]) &
                                (as.numeric(cell_floors) <= varRange[2])]

  # This function returns record counts up to, but not including,
  # varRange[2].
  n <- length(cell_floors_tmp) - 1
  counts <- rep(NA, n)
  for(i in 1:n) {
```

```

    lower <- as.numeric(cell_floors_tmp[i])
    upper <- as.numeric(cell_floors_tmp[i + 1])
    counts[i] <- length(med_houseVal[((med_houseVal >= lower) &
                                      (med_houseVal < upper))])
  }
  names(counts) <- names(cell_floors_tmp)[1:n]
  return(counts)
}

```

```

In [9]: cell_floors <- seq(from= 15000, to= 495000, by= 15000)
length(cell_floors)
names(cell_floors) <- paste(as.character(cell_floors/1000), "K", sep="")
print(cell_floors)

```

33

15K	30K	45K	60K	75K	90K	105K	120K	135K	150K	165K
15000	30000	45000	60000	75000	90000	105000	120000	135000	150000	165000
180K	195K	210K	225K	240K	255K	270K	285K	300K	315K	330K
180000	195000	210000	225000	240000	255000	270000	285000	300000	315000	330000
345K	360K	375K	390K	405K	420K	435K	450K	465K	480K	495K
345000	360000	375000	390000	405000	420000	435000	450000	465000	480000	495000

```

In [11]: observed_counts <- get_rcd_counts(dat$median_house_value, c(15000, 495000))
print(observed_counts)

```

15K	30K	45K	60K	75K	90K	105K	120K	135K	150K	165K	180K	195K	210K	225K	240K
14	78	587	965	1025	1295	1205	1088	1284	1538	1234	1239	842	931	959	747
255K	270K	285K	300K	315K	330K	345K	360K	375K	390K	405K	420K	435K	450K	465K	480K
656	615	439	355	353	400	386	266	207	190	163	135	115	134	81	68

```

In [12]: # Get the number of records not captured in observed_counts.

```

```
nrow(dat) - (sum(observed_counts) + 990)
```

19

```

In [13]: # The 19 records are between 495K and 500K.

```

```

nrow(dat[which((dat$median_house_value >= 495000) &
               (dat$median_house_value < 500000)),])

excluded_rows <- rownames(dat[which((dat$median_house_value >= 495000) &
                                   (dat$median_house_value < 500000)),])

```

19

```

In [14]: # Plot the counts. This will give us a very general idea
# of what the distribution of counts might look like for the
# 990 records which need an imputed value. We are especially
# interested in the general shape of the distribution from
# around 350K onwards.

```

```

df_plot <- rep(NA, 2 * length(observed_counts))
dim(df_plot) <- c(length(observed_counts), 2)
df_plot <- as.data.frame(df_plot)
colnames(df_plot) <- c("cell", "count")

new_names <- str_replace_all(names(observed_counts), "[K]", "")
df_plot$cell <- as.numeric(new_names)
df_plot$count <- as.numeric(observed_counts)

options(repr.plot.width= 13, repr.plot.height= 8)


p <- ggplot(df_plot, aes(cell, count)) +

```

```
geom_point(alpha= 0.5) + xlab("median_house_value floor") +
ylab("record count") +
geom_line() +
ggtitle("Figure 1: Count of records in each 15K bin of median_house_value") +
theme(axis.text= element_text(size= 12)) +
theme(axis.title= element_text(size= 14)) +
theme(title= element_text(size= 16))
```

p

Figure 1: Count of records in each 15K bin of median\_house\_value



In [15]: *# There is much less variability in the tail of the distribution.*

```
dim(df_plot)
print(sd(df_plot$count))
print(sd(df_plot[24:32,]$count))
```

32 · 2

[1] 460.4  
[1] 62.905

## Hypothesized distribution

We have a general idea of what the distribution of the actual, unobserved values will look like. Here I construct an example, or hypothesized, distribution. We know that as median house values increase, the number of districts (i.e., census blocks) will decrease; the correlation between these two variables is over 90%. Among the 990 records for which we need to impute a value, there are probably a few outliers, but we do not need to worry about trying to predict for these. Instead, we are interested in approximating what is likely to be the distribution for the vast majority of records.

The hypothesized distribution provides us with a way to judge the plausibility of our model-based prediction for the mean.

In [16]: *# Create an example, or hypothesized, distribution for the  
# expected range of imputation. (Previous work shows an  
# upper limit around 840K; so for this example distribution  
# I will go out only to the 825K bin.)*

```
bins <- seq(495000, 825000, by= 15000)
bin_names <- paste(as.character(bins/1000), "K", sep="")
```

```
names(bins) <- bin_names
names(bins)
length(bins)
# 23

# In addition to the 990 records to distribute, we have 19
# records that belong to the 495K cell.
bin_counts <- c(89, 95, 89, 86, 78, 71, 67, 65, 56, 51, 48,
                42, 38, 35, 27, 22, 17, 12, 7, 6, 2, 3, 3)
sum(bin_counts)
sum(bin_counts) == (990 + 19)
```

```
'495K'· '510K'· '525K'· '540K'· '555K'· '570K'· '585K'· '600K'· '615K'· '630K'· '645K'· '660K'· '675K'·
'690K'· '705K'· '720K'· '735K'· '750K'· '765K'· '780K'· '795K'· '810K'· '825K'
```

```
23
```

```
1009
```

```
TRUE
```

In [17]: *# Construct a dataframe for plotting of the example distribution.*

```
all_names <- c(df_plot$cell[24:32], bin_names)
observed <- df_plot$count[24:32]

all <- c(observed, bin_counts)
n <- length(all)

dftmp <- rep(NA, 2 * n)
dim(dftmp) <- c(n, 2)
dftmp <- as.data.frame(dftmp)
colnames(dftmp) <- c("cell", "count")
dftmp$cell <- all_names
dftmp$count <- all

dftmp$hhval <- as.numeric(str_replace_all(dftmp$cell, "[K]", ""))

head(dftmp); tail(dftmp)
```

A data.frame: 6 × 3

	cell	count	hhval
	<chr>	<dbl>	<dbl>
1	360	266	360
2	375	207	375
3	390	190	390
4	405	163	405
5	420	135	420
6	435	115	435


A data.frame: 6 × 3

	cell	count	hhval
	<chr>	<dbl>	<dbl>
27	750K	12	750
28	765K	7	765
29	780K	6	780
30	795K	2	795
31	810K	3	810
32	825K	3	825

```
In [18]: # Plot showing possible distribution of 990 + 19 districts  
# with a median_house_value >= 495K.
```

```
options(repr.plot.width= 11, repr.plot.height= 7)  
  
p <- ggplot(dftmp, aes(hhval, count)) +  
  geom_point(alpha= 0.5) + xlab("cell floor") + ylab("record count") +  
  geom_line() +  
  ggtitle("Figure 2: Possible distribution of counts >= 495K") +  
  theme(axis.text= element_text(size= 12)) +  
  theme(axis.title= element_text(size= 14)) +  
  theme(title= element_text(size= 16))  
p
```

Figure 2: Possible distribution of counts >= 495K



```

In [19]: # Compute the mean and median of our example distribution.
# These become our first estimates of the mean and median
# of the actual, unobserved median house values >= 500K.

dftmp <- dftmp[which(dftmp$hhval >= 495),]

# newvals will be used in cells downstream.
newvals <- c()
for(i in 1:nrow(dftmp)) {

  # Remove the 19 rcids >= 495K and < 500K.
  ifelse(i > 1, n <- dftmp$count[i], n <- dftmp$count[i] - 19)

  ifelse(i > 1, lower <- dftmp$hhval[i], lower <- dftmp$hhval[i] + 5)
  ifelse(i > 1, upper <- lower + 15, upper <- lower + 10)

  seed <- set.seed(4321 + i)
  vals <- round(runif(n, lower, upper))
  newvals <- c(newvals, vals)
}

length(newvals)
# 990
round(mean(newvals), 1)
# 599.7
round(median(newvals), 1)
# 586.5

```

990

599.7

586.5

```

In [ ]: ### COMMENTS:

# The example distribution has a mean of 600K. This is an
# estimate for the mean of the actual, unobserved median
# house values >= 500K. The estimate for the median is
# lower, as expected.

```

## Re-bin median\_house\_value using 10K intervals

```

In [20]: # We have 990 imputed values.

imputed_vals_tmp <- 1000*newvals

```

```

In [21]: # Combine the newly imputed values with the median house
# values in dat that are not censored.

all_hh_median_vals <- c(dat[which(dat$median_house_value < 500000), c("median_house_value")]
                        imputed_vals_tmp)
length(all_hh_median_vals)
summary(all_hh_median_vals)

```

20603

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
15000	119600	179800	211700	264950	836000

In [22]: *# Get record counts in each 10K bin from 50K up to 840K.*

```
all_counts <- get_rcd_counts(all_hh_median_vals, c(50000, 840000), span=10000, startpt=50000)
print(all_counts)
```

```
50K 60K 70K 80K 90K 100K 110K 120K 130K 140K 150K 160K 170K 180K 190K 200K
481 625 598 767 919 734 847 723 847 802 968 999 805 852 694 535
210K 220K 230K 240K 250K 260K 270K 280K 290K 300K 310K 320K 330K 340K 350K 360K
637 676 577 494 448 461 454 336 264 232 234 242 264 242 280 200
370K 380K 390K 400K 410K 420K 430K 440K 450K 460K 470K 480K 490K 500K 510K 520K
145 128 127 113 113 82 99 69 109 50 56 47 40 66 64 63
530K 540K 550K 560K 570K 580K 590K 600K 610K 620K 630K 640K 650K 660K 670K 680K
58 54 65 45 47 46 46 47 41 33 33 38 29 31 22 27
690K 700K 710K 720K 730K 740K 750K 760K 770K 780K 790K 800K 810K 820K 830K
25 18 18 16 11 14 9 5 5 3 3 2 2 3 1
```

In [23]: `length(all_counts)`  
`which(names(all_counts)== "500K")`

79

46

In [24]: *# The histogram below shows the counts for the example  
# distribution; this is a close-up of Figure 2.*

*## NOTE: In order for the following histograms to look right,  
## I need to add 1 to my values.*

```
tbl <- all_counts[46:79]
```

```
options(repr.plot.width= 15, repr.plot.height= 7)
```

```
mat <- t(as.matrix(c(1,2)))
layout(mat, widths = rep.int(20, ncol(mat)),
       heights = rep.int(7, nrow(mat)), respect = FALSE)
```

```
hist(rep(seq(500, 830, by=10), as.numeric(tbl)) + 1, breaks=30, xlab="median house value",
     main="Figure 3a: Example distribution for imputed values", ylim=c(0, 80))
```

```
hist(newvals + 1, breaks= 30, xlab="median house value", ylim=c(0, 80),
     main="Figure 3b: Example distribution with sampled imputed values")
```

Figure 3a: Example distribution for imputed values





Figure 3b: Example distribution with sampled imputed values



In [ ]: *### COMMENTS:*



```
# Following Appendix A, I rely on Figures 2, 3b, and 4 for
# judging the plausibility of predicted means and medians
# using the models that follow.
```

```
In [25]: # Smooth out the counts in our 10K bins.
```

```
print(tbl[1:30])
```

```
500K 510K 520K 530K 540K 550K 560K 570K 580K 590K 600K 610K 620K 630K 640K 650K
  66   64   63   58   54   65   45   47   46   46   47   41   33   33   38   29
660K 670K 680K 690K 700K 710K 720K 730K 740K 750K 760K 770K 780K 790K
  31   22   27   25   18   18   16   11   14    9    5    5    3    3
```

```
In [26]: tbl["500K"] <- 67; tbl["510K"] <- 65
tbl["530K"] <- 60; tbl["540K"] <- 57
tbl["550K"] <- 55; tbl["560K"] <- 51
```

```
tbl["570K"] <- 48; tbl["590K"] <- 43
tbl["600K"] <- 42; tbl["620K"] <- 38
tbl["630K"] <- 35; tbl["640K"] <- 32
tbl["670K"] <- 25
```

```
sum(tbl)
```


```
990
```

```
In [32]: # Check the shape of the revised hypothesized distribution.
```

```
options(repr.plot.width= 8, repr.plot.height= 6)
```

```
hist(rep(seq(500, 830, by=10), as.numeric(tbl)) + 1, breaks=30, xlab="median house value",
     main="Figure 4: Revised hypothesized distribution",
     cex.main= 1.6, ylim=c(0, 80))
```

**Figure 4: Revised hypothesized distribution**



```
In [28]: # Re-compute the mean and median of our hypothesized distribution.
```

```
# newvals will be used in cells downstream.
```

```
newvals <- c()
```

```
for(i in 1:length(tbl)) {
```

```
  n <- as.numeric(tbl[i])
  bin_name <- names(tbl[i])
  bin_name <- str_replace(bin_name, "K", "")
  lower <- as.numeric(bin_name)
  upper <- lower + 10
```

```

    seed <- set.seed(4321 + i)
    vals <- round(runif(n, lower, upper))
    newvals <- c(newvals, vals)
  }

length(newvals)
# 990
round(mean(newvals), 1)
# 600.1
round(median(newvals), 1)
# 586

```

990

600.1

586

```
In [208]: imputed_vals_tmp <- 1000 * newvals
```

```
In [ ]: ### COMMENTS:
```

```

# PREDICTION FOR THE MEAN: 600K
# PREDICTION FOR THE MEDIAN: 586K

```

## Compute shift-increment ratios for the mean with 300K window

I will start by using a rolling window of 300K. This window captures nearly all of the current example distribution of the imputed values when we start at the cap of 500K ( $500K + 300K = 800K$ ). Compute data for our prediction model from 50K - 330K. Although this takes us into the region of imputed values (we will use the example distribution of Figure 4), most of the data for the last few 300K windows will still be observed rather than imputed. See Appendix A for an example; by doing this, I should be able to obtain a far more accurate prediction for the mean. The hypothesized distribution shows us what we think will happen between 500K and 840K; revisit Figure 2 to judge the plausibility in the larger context. Another plausible distribution is one in which the counts drop more quickly and in which the tail extends much further to the right, and/or we have a few distant outliers. While distant outliers can change the mean quite a bit, I am interested in only in trying to get good predictions for the vast majority of the median house values  $\geq 500K$ . Thus our focus should be on where the mean might lie for this majority of values, nearly all of which will lie below 800K. (Keep in mind that the median is for a census block; on average these blocks have around 400 households. This makes it difficult for the median values to get too high.)

In Appendix A we saw that predicting the median is not so easy. Here I am not going to worry about getting a prediction for the median. As in Appendix A, we know that the median will lie somewhere to the left of the mean, and our hypothesized distribution gives us some idea of the expected distance between the mean and the median.

We will see below that the 300K window is too large; we are looking out further than we ought to for bins  $< 500K$ . This inflates the ratios we rely on for our prediction; we thus end up with a prediction that is much higher than the 600K prediction we have from the hypothesized distribution. The solution is to reduce the size of the window, perhaps down to 180K.

```
In [33]: bins <- seq(50000, 330000, by= 10000)
bin_names <- paste(as.character(bins/1000), "K", sep="")
names(bins) <- bin_names
length(bins)
```

29

```
In [209]: # See Figure 3b.
summary(newvals)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
500	539	586	600	650	838

```
In [210]: # Combine the newly imputed values with the median house
```

```
# values in dat that are not censored.

all_hh_median_vals <- c(dat[which(dat$median_house_value < 500000), c("median_house_value")]
                        newvals*1000)
length(all_hh_median_vals)
summary(all_hh_median_vals)

20603

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 15000  119600  179800  211717  264950  838000
```

```
In [36]: # Get the means for each bin, using a 300K window. Note that 300K
# is divisible by 10K, the size of each median_house_value bin.
# (This is important because it means that we are never breaking
# a bin apart when calling get_rcd_counts in the loop below.)

mean_ratios <- rep(NA, length(bins))
means <- rep(NA, length(bins))
rcd_count <- rep(NA, length(bins))

span <- 300000
index <- 0
for(floor in bins) {

  index <- index + 1
  hhvals <- as.numeric(all_hh_median_vals[which((all_hh_median_vals >= floor) &
                                                (all_hh_median_vals < (floor + span))]))
  counts <- as.numeric(get_rcd_counts(hhvals, c(floor, (floor+span)),
                                       span=10000, startpt=50000))
  rcd_count[index] <- sum(counts)

  # Compute mean.
  hhval_mean <- round(mean(hhvals), 5)
  mean_ratios[index] <- round(hhval_mean/floor, 3)
  means[index] <- hhval_mean

}

paste0("These are the 300K shift increments for the means: ")
names(mean_ratios) <- bin_names
print(mean_ratios)
```

'These are the 300K shift increments for the means: '

50K	60K	70K	80K	90K	100K	110K	120K	130K	140K	150K	160K	170K
3.481	3.003	2.665	2.402	2.209	2.067	1.940	1.842	1.750	1.682	1.619	1.580	1.545
180K	190K	200K	210K	220K	230K	240K	250K	260K	270K	280K	290K	300K
1.509	1.482	1.452	1.422	1.403	1.392	1.380	1.368	1.356	1.350	1.348	1.340	1.328
310K	320K	330K										
1.316	1.306	1.299										


```
In [37]: # Construct dataframe for plotting, etc.

df_ratios <- rep(NA, 4*length(mean_ratios))
dim(df_ratios) <- c(length(mean_ratios), 4)
df_ratios <- as.data.frame(df_ratios)
colnames(df_ratios) <- c("cell", "rcds", "mean", "mean_ratio")
df_ratios$cell <- bins
df_ratios$rcds <- rcd_count
df_ratios$mean_ratio <- mean_ratios
df_ratios$mean <- means
```

```
In [38]: options(repr.plot.width= 12, repr.plot.height= 8)

p <- ggplot(df_ratios, aes(cell, mean_ratio)) +
  geom_point(alpha= 0.5) + xlab("median house value") +
  ylab("300K ratio: means") +
  geom_line() +
```

```
ggtitle("300K shift increment ratios for means") +
theme(axis.text= element_text(size = 12)) +
theme(axis.title= element_text(size= 14)) +
theme(title= element_text(size= 16))
p
```



```
In [39]: # We cannot bring rcdds into our model because it is too
# highly correlated with cell (median_house_value). We
# might need weights on cell to establish constant variance.

df_ratios$rcdds
```

```
17757 · 17556 · 17131 · 16678 · 16039 · 15247 · 14626 · 13892 · 13251 · 12503 · 11770 · 10911 ·
9962 · 9213 · 8408 · 7754 · 7282 · 6709 · 6099 · 5581 · 5143 · 4747 · 4340 · 3935 · 3646 · 3424 ·
3235 · 3044 · 2838
```

```
In [34]: # Compute correlation between rcdds and cell. rcdds go down
# in count as cell (median house value) increases. If
# these predictors are too highly correlated, I should
# remove rcdds from the model.

round(cor(df_ratios$rcdds^0.15, df_ratios$cell^-0.385), 3)

0.91
```

### Try weighted least squares

```
In [144]: # Use rcdds as a surrogate for cell in the weights.

g04 <- lm(I(mean_ratio^0.25) ~ I(cell^0.025) + I((cell^0.025)^2),
          data= df_ratios, weights=df_ratios$rcdds^0.025)

ans <- summary(g04)
ans[[1]] <- ""; ans
```

Call:

""

Weighted Residuals:

	Min	1Q	Median	3Q	Max
	-0.003143	-0.001195	0.000200	0.000906	0.002744

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	128.22	1.45	88.6	<2e-16
I(cell <sup>0.025</sup> )	-184.65	2.15	-85.8	<2e-16
I((cell <sup>0.025</sup> ) <sup>2</sup> )	67.04	0.80	83.8	<2e-16

Residual standard error: 0.00149 on 26 degrees of freedom

In [145]: `ncvTest(g04)`


Non-constant Variance Score Test  
Variance formula: ~ fitted.values  
Chisquare = 1.4571, Df = 1, p = 0.227

In [146]: `residualPlots(g04, plot=FALSE)`

	Test stat	Pr(> Test stat )
I(cell <sup>0.025</sup> )	3.11	0.0047
I((cell <sup>0.025</sup> ) <sup>2</sup> )	2.60	0.0153
Tukey test	0.01	0.9924

In [147]: `options(repr.plot.width= 6, repr.plot.height= 6)`

```
ans <- qqnorm(scale(residuals(g04, type= "pearson")))
qqline(ans$x, probs = c(0.25, 0.75))
```



```
In [148]: # Prediction for mean for [500K, 800K].

newdat <- df_ratios[1, ]
newdat[1, ] <- c(500000, 990, rep(NA, 2))

ans <- predict.lm(g04, newdata= newdat, type= "response")
ans_transf <- ans^(1/0.25); ans_transf
# 1.35

# 1.35 * 500K = 675K.
```

```
1: 1.34995422675475
```

```
In [131]: # Compute a 95% prediction interval.

pred_ans <- predict.lm(g04, newdata= newdat, interval="prediction",
                      level=0.95, weights=990^0.025)
pred_ans_transf <- pred_ans^(1/0.25); pred_ans_transf
```

```
A matrix: 1 x 3 of type dbl
```

	fit	lwr	upr
1	1.35	1.3304	1.3697

```
In [132]: lwr <- round(pred_ans_transf[2] * 500)
upr <- round(pred_ans_transf[3] * 500)

clause <- "95% prediction interval for estimate of the mean of the actual, unobserved values"
print_ans <- paste0("[", lwr, "K, ", upr, "K]")
paste0(clause, print_ans)
# [665K, 685K]
```

```
'95% prediction interval for estimate of the mean of the actual, unobserved values: [665K, 685K]'
```

```
In [236]: ### COMMENT:

# The prediction interval does not come close to including our
# earlier prediction of 600K. And thus, we can conclude that
# this model prediction is of no help to us. The problem
# lies with the window size. Once we reduce it, we will get
# a prediction that lies closer to 600K.
```

## Return to ordinary least squares

```
In [154]: # The weights are not making much of a difference in g04.
# So the g05 model does not make use of them.

g05 <- lm(I(mean_ratio^0.24) ~ I(cell^0.025) + I((cell^0.025)^2),
        data= df_ratios)

ans <- summary(g05)
ans[[1]] <- ""; ans
```

Call:  
""

Residuals:  
Min 1Q Median 3Q Max

In [155]: `ncvTest(g05)`


Non-constant Variance Score Test  
Variance formula:  $\sim$  fitted.values  
Chisquare = 1.2072, Df = 1, p = 0.272

In [156]: `residualPlots(g05, plot=FALSE)`

	Test stat	Pr(> Test stat )
I(cell <sup>0.025</sup> )	0.59	0.56
I((cell <sup>0.025</sup> ) <sup>2</sup> )	0.10	0.92
Tukey test	-0.48	0.63

In [157]: `options(repr.plot.width= 6, repr.plot.height= 6)`

```
ans <- qqnorm(scale(residuals(g05, type= "pearson")))
qqline(ans$x, probs = c(0.25, 0.75))
```



In [158]: *# Prediction for mean for [500K, 800K].*

```
newdat <- df_ratios[1, ]
newdat[1, ] <- c(500000, 990, rep(NA, 2))

ans <- predict.lm(g05, newdata= newdat, type= "response")
ans_transf <- ans^(1/0.24); ans_transf
# 1.3477

# 1.3477 * 500K = 674K.
```

1: 1.34771224418045

In [159]: *# Compute a 95% prediction interval.*

```
pred_ans <- predict.lm(g05, newdata= newdat, interval="prediction",
  level=0.95)
pred_ans_transf <- pred_ans^(1/0.24); pred_ans_transf
```

A matrix: 1 × 3 of type dbl

fit      lwr      upr

```
In [160]: lwr <- round(pred_ans_transf[2] * 500)
upr <- round(pred_ans_transf[3] * 500)

clause <- "95% prediction interval for estimate of the mean of the actual, unobserved values
print_ans <- paste0("[", lwr, "K, ", upr, "K]")
paste0(clause, print_ans)
# [664K, 684K]
```

'95% prediction interval for estimate of the mean of the actual, unobserved values: [664K, 684K]'

## Compute shift-increments using a 220K window

I can size the window by looking at the last ratio in the sequence I rely on for my prediction. If that ratio is < 1.2, then I know the window is too small. For  $1.2 * 500K = 600K$ , and 600K is the prediction we have from the hypothesized distribution. So our last ratio should be > 1.2 since we are still far to the left of 500K. When I tried a window of size 180K, the last ratio (for 370K) was 1.196 (and at 330K the ratio was even lower, at 1.186). So here I will try 220K.

```
In [42]: bins <- seq(50000, 340000, by= 10000)
bin_names <- paste(as.character(bins/1000), "K", sep="")
names(bins) <- bin_names
length(bins)
```

30

```
In [43]: mean_ratios <- rep(NA, length(bins))
means <- rep(NA, length(bins))
rcd_count <- rep(NA, length(bins))

span <- 220000
index <- 0
for(floor in bins) {
  index <- index + 1
  hhvals <- as.numeric(all_hh_median_vals[which((all_hh_median_vals >= floor) &
                                                (all_hh_median_vals < (floor + span)))]])
  counts <- as.numeric(get_rcd_counts(hhvals, c(floor, (floor+span)),
                                       span=10000, startpt=50000))
  rcd_count[index] <- sum(counts)

  # Compute mean.
  hhval_mean <- round(mean(hhvals), 5)
  mean_ratios[index] <- round(hhval_mean/floor, 3)
  means[index] <- hhval_mean
}

paste0("These are the 300K shift increments for the means: ")
names(mean_ratios) <- bin_names
print(mean_ratios)
```

'These are the 300K shift increments for the means: '

50K	60K	70K	80K	90K	100K	110K	120K	130K	140K	150K	160K	170K
3.097	2.691	2.402	2.177	2.012	1.892	1.785	1.707	1.634	1.582	1.532	1.494	1.466
180K	190K	200K	210K	220K	230K	240K	250K	260K	270K	280K	290K	300K
1.436	1.413	1.390	1.360	1.342	1.328	1.319	1.303	1.289	1.278	1.272	1.264	1.253
310K	320K	330K	340K									
1.243	1.234	1.228	1.226									


```
In [44]: # Construct dataframe for plotting, etc.
```



```
df_ratios <- rep(NA, 4*length(mean_ratios))
dim(df_ratios) <- c(length(mean_ratios), 4)
df_ratios <- as.data.frame(df_ratios)
colnames(df_ratios) <- c("cell", "rcds", "mean", "mean_ratio")
df_ratios$cell <- bins
df_ratios$rcds <- rcd_count
df_ratios$mean_ratio <- mean_ratios
df_ratios$mean <- means
```

```
In [45]: options(repr.plot.width= 12, repr.plot.height= 8)

p <- ggplot(df_ratios, aes(cell, mean_ratio)) +
  geom_point(alpha= 0.5) + xlab("median house value") +
  ylab("220K ratio: means") +
  geom_line() +
  ggtitle("220K shift increment ratios for means") +
  theme(axis.text= element_text(size = 12)) +
  theme(axis.title= element_text(size= 14)) +
  theme(title= element_text(size= 16))
p
```



```
In [96]: g06 <- lm(I(mean_ratio^0.9) ~ I(cell^0.02)
  + I((cell^0.02)^2) + I((cell^0.02)^3),
  data= df_ratios)

ans <- summary(g06)
ans[[1]] <- ""; ans
```

Call:  
""

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

In [97]: `ncvTest(g06)`


Non-constant Variance Score Test  
Variance formula:  $\sim$  fitted.values  
Chisquare = 0.0024245, Df = 1, p = 0.961

In [98]: `residualPlots(g06, plot=FALSE)`

	Test stat	Pr(> Test stat )
I(cell <sup>0.02</sup> )	-0.26	0.80
I((cell <sup>0.02</sup> ) <sup>2</sup> )	0.45	0.66
I((cell <sup>0.02</sup> ) <sup>3</sup> )	0.46	0.65
Tukey test	0.08	0.93

In [99]: `options(repr.plot.width= 6, repr.plot.height= 6)`

```
ans <- qqnorm(scale(residuals(g06, type= "pearson")))
qqline(ans$x, probs = c(0.25, 0.75))
```



In [100]: *# Prediction for mean for [500K, 800K].*

```
newdat <- df_ratios[1, ]
newdat[1, ] <- c(500000, 990, rep(NA, 2))

ans <- predict.lm(g06, newdata= newdat, type= "response")
ans_transf <- ans^(1/0.9); ans_transf
# 1.1138

# 1.1138 * 500K = 557K.
```

1: 1.1137883449534

In [101]: *# Compute a 95% prediction interval.*

```
pred_ans <- predict.lm(g06, newdata= newdat, interval="prediction",
                      level=0.95)
pred_ans_transf <- pred_ans^(1/0.9); pred_ans_transf
```

A matrix: 1 × 3 of type dbl

	fit	lwr	upr
1	1.1138	1.0966	1.131

```
In [102]: lwr <- round(pred_ans_transf[2] * 500)
upr <- round(pred_ans_transf[3] * 500)

clause <- "95% prediction interval for estimate of the mean of the actual, unobserved values
print_ans <- paste0("[", lwr, "K, ", upr, "K]")
paste0(clause, print_ans)
# [548K, 566K]
```

'95% prediction interval for estimate of the mean of the actual, unobserved values: [548K, 566K]'

```
In [ ]: ### COMMENT:

# The 95% prediction interval should cover the 600K prediction.
# Since it does not when our window is at 220K, we should
# increase the window size. Increasing the size of our rolling
# window increases the shift-increment ratios (a surrogate for
# the mean).
```

## Compute shift-increments using a 240K window

If we use a 250K window, our 500K point estimate is 626K. This is too high, since the 95% prediction interval will not cover the 600K prediction we already have. So here I am trying a window that spans 240K.

```
In [160]: # By stopping at 330K, the last ratio relies on data
# out to 570K. As usual, we rely somewhat on the
# hypothesized distribution.

bins <- seq(50000, 330000, by= 10000)
bin_names <- paste(as.character(bins/1000), "K", sep="")
names(bins) <- bin_names
length(bins)
```

29

```

In [161]: mean_ratios <- rep(NA, length(bins))
          means <- rep(NA, length(bins))
          rcd_count <- rep(NA, length(bins))

          span <- 240000
          index <- 0
          for(floor in bins) {

            index <- index + 1
            hhvals <- as.numeric(all_hh_median_vals[which((all_hh_median_vals >= floor) &
                                                         (all_hh_median_vals < (floor + span))]])
            counts <- as.numeric(get_rcd_counts(hhvals, c(floor, (floor+span)),
                                                         span=10000, startpt=50000))
            rcd_count[index] <- sum(counts)

            # Compute mean.
            hhval_mean <- round(mean(hhvals), 5)
            mean_ratios[index] <- round(hhval_mean/floor, 3)
            means[index] <- hhval_mean

          }

          paste0("These are the 300K shift increments for the means: ")
          names(mean_ratios) <- bin_names
          print(mean_ratios)

```

'These are the 300K shift increments for the means: '

50K	60K	70K	80K	90K	100K	110K	120K	130K	140K	150K	160K	170K
3.217	2.770	2.461	2.229	2.061	1.942	1.833	1.755	1.677	1.613	1.556	1.517	1.488
180K	190K	200K	210K	220K	230K	240K	250K	260K	270K	280K	290K	300K
1.457	1.433	1.409	1.377	1.361	1.346	1.332	1.316	1.300	1.293	1.291	1.284	1.274
310K	320K	330K										
1.262	1.253	1.247										

```


In [162]: # Construct dataframe for plotting, etc.

          df_ratios <- rep(NA, 4*length(mean_ratios))
          dim(df_ratios) <- c(length(mean_ratios), 4)
          df_ratios <- as.data.frame(df_ratios)
          colnames(df_ratios) <- c("cell", "rcds", "mean", "mean_ratio")
          df_ratios$cell <- bins
          df_ratios$rcds <- rcd_count
          df_ratios$mean_ratio <- mean_ratios
          df_ratios$mean <- means

```

```
In [163]: options(repr.plot.width= 12, repr.plot.height= 8)

p <- ggplot(df_ratios, aes(cell, mean_ratio)) +
  geom_point(alpha= 0.5) + xlab("median house value") +
  ylab("240K ratio: means") +
  geom_line() +
  ggtitle("240K shift increment ratios for means") +
  theme(axis.text= element_text(size= 12)) +
  theme(axis.title= element_text(size= 14)) +
  theme(title= element_text(size= 16))
p
```



```
In [178]: g07 <- lm(I(mean_ratio^0.27) ~ I(cell^0.04)
+ I((cell^0.04)^2) + I((cell^0.04)^3),
data= df_ratios)

ans <- summary(g07)
ans[[1]] <- ""; ans
```

Call:

```
lm(I(mean_ratio^0.27) ~ I(cell^0.04) + I((cell^0.04)^2) + I((cell^0.04)^3), data=df_ratios)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.96e-03	-7.01e-04	2.03e-05	4.64e-04	1.78e-03

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	228.04	21.40	10.66	8.8e-11
I(cell^0.04)	-390.24	40.02	-9.75	5.3e-10
I((cell^0.04)^2)	223.46	24.94	8.96	2.8e-09
I((cell^0.04)^3)	-42.62	5.18	-8.23	1.4e-08

Residual standard error: 0.001 on 25 degrees of freedom  
Multiple R-squared: 1, Adjusted R-squared: 1  
F-statistic: 6.17e+04 on 3 and 25 DF, p-value: <2e-16

```
In [179]: ncvTest(g07)
```


```
Non-constant Variance Score Test
Variance formula: ~ fitted.values
Chisquare = 0.42315, Df = 1, p = 0.515
```

```
In [180]: residualPlots(g07, plot=FALSE)
```

	Test stat	Pr(> Test stat )
I(cell <sup>0.04</sup> )	0.07	0.95
I((cell <sup>0.04</sup> ) <sup>2</sup> )	-0.08	0.94
I((cell <sup>0.04</sup> ) <sup>3</sup> )	-0.09	0.93
Tukey test	0.01	1.00

```
In [181]: options(repr.plot.width= 6, repr.plot.height= 6)

ans <- qqnorm(scale(residuals(g07, type= "pearson")))
qqline(ans$x, probs = c(0.25, 0.75))
```



```
In [182]: # Prediction for mean for [500K, 800K].

newdat <- df_ratios[1, ]
newdat[1, ] <- c(500000, 990, rep(NA, 2))

ans <- predict.lm(g07, newdata= newdat, type= "response")
ans_transf <- ans^(1/0.27); ans_transf
# 1.2093

# 1.2093 * 500K = 604.65K.

1: 1.20934467746617
```

```
In [183]: # Compute a 95% prediction interval.

pred_ans <- predict.lm(g07, newdata= newdat, interval="prediction",
                      level=0.95)
pred_ans_transf <- pred_ans^(1/0.27); pred_ans_transf
```

A matrix: 1 × 3 of type dbl

	fit	lwr	upr
1	1.2093	1.1864	1.2326

```
In [184]: lwr <- round(pred_ans_transf[2] * 500)
upr <- round(pred_ans_transf[3] * 500)
```

```

clause <- "95% prediction interval for estimate of the mean of the actual, unobserved values
print_ans <- paste0("[", lwr, "K, ", upr, "K]")
paste0(clause, print_ans)
# [593K, 616K]

```

'95% prediction interval for estimate of the mean of the actual, unobserved values: [593K, 616K]'

## Final Comments for Section 1

As noted in Appendix A, if we have a good prediction for the mean of the actual, unobserved values (by "good" I mean a prediction which we can have a degree of confidence in), then we can improve upon the Gibbs sampler output if its predictions have a mean where we do not think it actually is. From the above we can be fairly confident that the mean for the median house values  $\geq 500K$  is close to **605K**. We can be fairly confident that the median will be less than the mean.

As in Appendix A, it is much harder to predict for the median than it is for the mean.

We have seen in the above analysis that the size of the shift-increment window has a significant effect on the prediction we get. We can determine an appropriate size for the window by relying on our prediction for the mean based on the hypothesized distribution. If in the above we set the window to 220K, the 95% prediction interval is too low to include the 600K estimate. If we set the window to 250K, the 95% prediction interval is too high to include the 600K estimate. When we use a 240K window, however, we get a very plausible prediction---again, assuming that our hypothesized distribution provides us with what we take to be a reasonable estimate.

In [ ]:

## Section 2: impute values for censored median house values

```

In [185]: # The following model is what we will use to predict the
# median house values that we need.

m01 <- lm(I(median_house_value^0.18) ~

      I(median_income^0.77) +
      I(long_transf^-0.5) +
      I(long_transf^-1) +
      I(long_transf^-1.5) +
      latitude +
      I(latitude^2) +
      I(latitude^3) +
      I(latitude^4) +
      pop_per_hh +
      I(pop_per_hh^2) +
      I(housing_median_age^0.15) +
      HHdens_ln +
      HHdens_ln:long_transf +
      HHdens_ln:median_income +
      HHdens_ln:housing_median_age:median_income,

      data= dat)

# REMARK: dat includes the capped data; if we discard the censored
# records, we are discarding valuable information.

m01.summary <- summary(m01)
m01.summary[[1]] <- ""; round(m01.summary$adj.r.squared, 3)

```

0.73

In [186]: `ncvTest(m01)`


Non-constant Variance Score Test  
 Variance formula: ~ fitted.values

```
In [187]: residualPlots(m01, plot=FALSE)
```

	Test stat	Pr(> Test stat )
I(median_income^0.77)	-14.13	<2e-16
I(long_transf^-0.5)	1.99	0.046
I(long_transf^-1)	11.11	<2e-16
I(long_transf^-1.5)	11.55	<2e-16
latitude	0.89	0.373
I(latitude^2)	-0.40	0.692
I(latitude^3)	33.30	<2e-16
I(latitude^4)	33.28	<2e-16
pop_per_hh	-1.32	0.186
I(pop_per_hh^2)	-13.36	<2e-16
I(housing_median_age^0.15)	0.46	0.643
HHdens_ln	11.34	<2e-16
Tukey test	0.07	0.945

```
In [188]: options(repr.plot.width= 6, repr.plot.height= 6)

ans <- qqnorm(scale(residuals(m01, type= "pearson")))
qqline(ans$x, probs = c(0.25, 0.75))
```



```
In [189]: # Get a sense of the uncertainty for the model's sigma.
# (sim is from the arm package.)

m01.sim <- sim(m01, n.sims=3000)
```


```
In [190]: sigma.m01.sim <- sigma.hat(m01.sim)
str(sigma.m01.sim)

num [1:3000] 0.46 0.469 0.47 0.465 0.465 ...
```



```
In [191]: options(repr.plot.width= 8, repr.plot.height= 6)

hist(sigma.m01.sim, breaks=20, main="Distribution of m01's sigma",
      xlab="model m01's sigma distribution")
```



```
In [ ]: # sigma.hat is small because of the power transformation
# on the response variable.
```

### Gibbs sampler for imputing censored median\_house\_values

```
In [192]: # Because of the transformation on the response variable,
# we need to transform our limits. Here I am setting the
# upper limit to 800K.

cap <- 500000
response_var_power <- 0.18
inv_pwr <- 1/response_var_power

# Set C_upper to where we think we will have accounted for
# a vast majority of the actual, unobserved values.
C <- cap^response_var_power
C_upper <- (1.6*cap)^response_var_power

censored <- (dat$median_house_value)^response_var_power >= C

# Create some crude starting values.
n.censored <- sum(censored)
z <- ifelse(censored, NA, (dat$median_house_value)^response_var_power)
z[censored] <- runif(n.censored, C, C_upper)
```

```
In [193]: length(censored)
n.censored
```

20603

990

```
In [194]: summary(z[censored])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
10.6	10.8	11.1	11.1	11.3	11.5

In [195]: *# Identify the rows that are censored.*

```
rows_censored <- rownames(dat)[censored]
head(rows_censored)
```

```
'90' · '460' · '494' · '495' · '510' · '511'
```

In [196]: *# Function to draw from a constrained normal distribution.*

```
rnorm.trunc03 <- function(n, mu, sigma, lo=-Inf, hi=Inf) {
  # We need each mu to be >= C. Otherwise the return
  # value will be Inf.
  cap <- 500000
  mu02 <- ifelse(mu <= C, (cap + 100)^response_var_power, mu)

  p.lo <- pnorm(lo, mu02, sigma)
  p.hi <- pnorm(hi, mu02, sigma)
  u <- runif(n, p.lo, p.hi)
  return(qnorm(u, mu02, sigma))
}
```

In [246]: *# Create matrix X for the terms in our model.*

```
X <- dat
X$median_income <- (X$median_income)^0.77

X$lat2 <- (X$latitude)^2
X$lat3 <- (X$latitude)^3
X$lat4 <- (X$latitude)^4

X$long_1 <- (X$long_transf)^-0.5
X$long_2 <- (X$long_transf)^-1
X$long_3 <- (X$long_transf)^-1.5

X$pphh1 <- X$pop_per_hh
X$pphh2 <- (X$pop_per_hh)^2

X$housing_median_age <- (X$housing_median_age)^0.15

X$HHdens_by_long <- X$HHdens_ln * X$long_transf
X$HHdens_by_income <- X$HHdens_ln * X$median_income
X$HHdens_3way <- X$HHdens_ln * X$median_income * X$housing_median_age

X <- X[, c("median_income", "long_1", "long_2", "long_3", "latitude", "lat2",
           "lat3", "lat4", "pphh1", "pphh2", "housing_median_age",
           "HHdens_ln", "HHdens_by_long", "HHdens_by_income",
           "HHdens_3way")]
intercept <- rep(1, nrow(dat))

init.colnames <- colnames(X)

X <- as.data.frame(cbind(intercept, X), col.names=c("intercept", init.colnames),
                  row.names=rownames(dat))
dim(X)
colnames(X)
```

```
20603 · 16
```

```
'intercept' · 'median_income' · 'long_1' · 'long_2' · 'long_3' · 'latitude' · 'lat2' · 'lat3' · 'lat4' · 'pphh1' · 'pphh2' ·
'housing_median_age' · 'HHdens_ln' · 'HHdens_by_long' · 'HHdens_by_income' · 'HHdens_3way'
```

In [247]: *# See p.406 (Section 18.5) of Gelman and Hill's book,*

```

# "Data Analysis Using Regression and Multilevel/Hierarchical
# Models".

# Fit a regression using the crude starting values of z.

m01_tst <- lm(z ~

      I(median_income^0.77) +
      I(long_transf^-0.5) +
      I(long_transf^-1) +
      I(long_transf^-1.5) +
      latitude +
      I(latitude^2) +
      I(latitude^3) +
      I(latitude^4) +
      pop_per_hh +
      I(pop_per_hh^2) +
      I(housing_median_age^0.15) +
      HHdens_ln +
      HHdens_ln:long_transf +
      HHdens_ln:median_income +
      HHdens_ln:housing_median_age:median_income,

      data= dat)

# Obtain a sample draw of the model coefficients and of
# parameter sigma.
sim.1 <- sim(m01_tst, n.sims=1)

```

```

In [248]: beta <- coef(sim.1)
          dim(beta)
          colnames(beta)

```

```
1 · 16
```

```

'(Intercept)' · 'I(median_income^0.77)' · 'I(long_transf^-0.5)' · 'I(long_transf^-1)' · 'I(long_transf^-1.5)' · 'latitude' ·
'I(latitude^2)' · 'I(latitude^3)' · 'I(latitude^4)' · 'pop_per_hh' · 'I(pop_per_hh^2)' · 'I(housing_median_age^0.15)' ·
'HHdens_ln' · 'HHdens_ln:long_transf' · 'HHdens_ln:median_income' ·
'HHdens_ln:median_income:housing_median_age'

```

```

In [249]: # Here are means for 6 different normal
          # distributions.

```

```

means <- as.matrix(X) %*% t(beta)
length(means)
round(head(as.vector(means)^inv_pwr))

```

```
20603
```

```
466164 · 520490 · 366964 · 297426 · 230703 · 241224
```

```

In [250]: # All values should be between 500K and 800K
          z.old <- z[censored]
          round(head(z.old)^inv_pwr)

```

```
507102 · 504066 · 558000 · 592871 · 799517 · 620749
```

```

In [251]: # All values should be between 500K and 800K.

```

```

sigma <- sigma.hat(sim.1)
round(sigma, 4)

z.new <- rnorm.trunc03(n.censored, means[censored], sigma, lo=C, hi=C_upper)
round(head(as.vector(z.new)^inv_pwr))

```

```
0.4976
```

```
574649 · 709798 · 616496 · 649026 · 622094 · 656843
```

In [252]: `summary(z.new^inv_pwr)`

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
500040  558565  621678  628473  695964  799530

```

In [253]: *# For the Gibbs sampler, the above is now put into  
# a loop. We first test with 100 iterations.*

```

n <- nrow(dat)
n.chains <- 4
n.iter <- 2000

sims <- array(NA, c(n.iter, n.chains, 17 + n.censored))
dimnames(sims) <- list(NULL, NULL, c(colnames(X), "sigma",
                                     paste("z[", (1:n)[censored],
                                             "]", sep="")))

start <- Sys.time()
for(m in 1:n.chains) {

  # acquire some initial values
  z[censored] <- runif(n.censored, C, C_upper)

  for(t in 1:n.iter) {

    m01.1 <- lm(z ~

      I(median_income^0.77) +
      I(long_transf^-0.5) +
      I(long_transf^-1) +
      I(long_transf^-1.5) +
      latitude +
      I(latitude^2) +
      I(latitude^3) +
      I(latitude^4) +
      pop_per_hh +
      I(pop_per_hh^2) +
      I(housing_median_age^0.15) +
      HHdens_ln +
      HHdens_ln:long_transf +
      HHdens_ln:median_income +
      HHdens_ln:housing_median_age:median_income,

      data= dat)

    sim.1 <- sim(m01.1, n.sims=1)
    beta <- coef(sim.1)
    sigma <- sigma.hat(sim.1)
    means <- as.matrix(X) %*% t(beta)
    z[censored] <- rnorm.trunc03(n.censored, means[censored], sigma, lo=C, hi=C_upper)
    stopifnot(sum(z[censored] < Inf) == n.censored)
    sims[t,m,] <- c(beta, sigma, z[censored])
  }
}

stop <- Sys.time()
round(stop - start, 2)
# Time difference of 3.88 minutes.

```

Time difference of 3.88 mins

In [242]: *# Check for convergence.*

```

# sims.bugs <- R2OpenBUGS::as.bugs.array(sims, n.burnin=1000)
# print(sims.bugs)

# The Rhat value for every parameter and every imputed
# value should be 1.0.

```

```
In [254]: save(sims, file="/home/greg/Documents/stat/Geron_ML/datasets/housing/sims_raw_hhvals.RData")
```

```
In [197]: load("/home/greg/Documents/stat/Geron_ML/datasets/housing/sims_raw_hhvals.RData")
```

```
In [198]: # Drop the first 1000 iterations.
```

```
sims_adj <- sims[1001:2000, ,]
dim(sims_adj)
```

```
1000 4 1007
```

```
In [199]: sims_adj.bugs <- R2OpenBUGS::as.bugs.array(sims_adj)
# print(sims_adj.bugs)
```

```
In [200]: # Extract the means and stddevs for each of the censored records.
```

```
z_means <- sims_adj.bugs$mean$z
z_sds <- sims_adj.bugs$sd$z
round(head(z_means), 2); round(head(z_sds), 2)
```

```
10.96 10.95 10.96 10.96 10.96 11.17
```

```
0.23 0.24 0.24 0.24 0.24 0.24
```

```
In [201]: summary(z_means)
summary(z_sds)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
10.9	11.0	11.0	11.0	11.1	11.4
Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.103	0.236	0.239	0.237	0.245	0.257

```
In [202]: summary(round(z_means^inv_pwr))
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
594362	597354	598802	620628	633302	758891

```
In [203]: # Average estimate of the sd.
```

```
(sd_estimate <- round((11 + 0.237)^inv_pwr) - round(11^inv_pwr))
# 76,724
```

```
76724
```

```
In [204]: # Here is a fuller summary for the stddevs.
```

```
ans <- round((z_means + z_sds)^inv_pwr) - round(z_means^inv_pwr)
summary(ans)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
38816	75314	76235	77571	81054	87973

```
In [ ]: ### COMMENTS:
```

```
# Based on the work above, we expect the mean to be 600K -
# 605K if the upper limit is around 770K. The mean is
# currently around 629K (see next summary).
```

```
In [205]: # Get some predictions, using rnorm.trunc03.
```

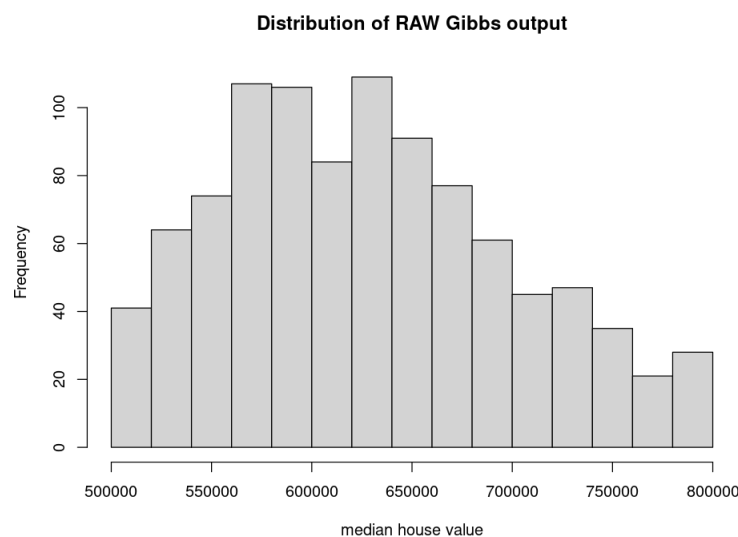
```
set.seed(1931)
z_preds <- round(rnorm.trunc03(n.censored, z_means, z_sds, lo=C, hi=C_upper), 5)
z_preds <- round(z_preds^inv_pwr)
summary(z_preds)
```

```
# Notice that the mean is at 629K. We do not expect the mean
# to be this high. In fact, this mean is not even in the
# 95% prediction interval from the g07 model prediction. The
# upper limit for that prediction interval is 616K.
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
500782  573851  624113  629122  677362  799914
```

```
In [206]: options(repr.plot.width= 8, repr.plot.height= 6)
```

```
hist(z_preds + 1, breaks=20, main="Distribution of RAW Gibbs output",
     xlab="median house value")
```



```
In [211]: # Compare hypothesized distribution to the above distribution.
```

```
options(repr.plot.width= 15, repr.plot.height= 6)
```

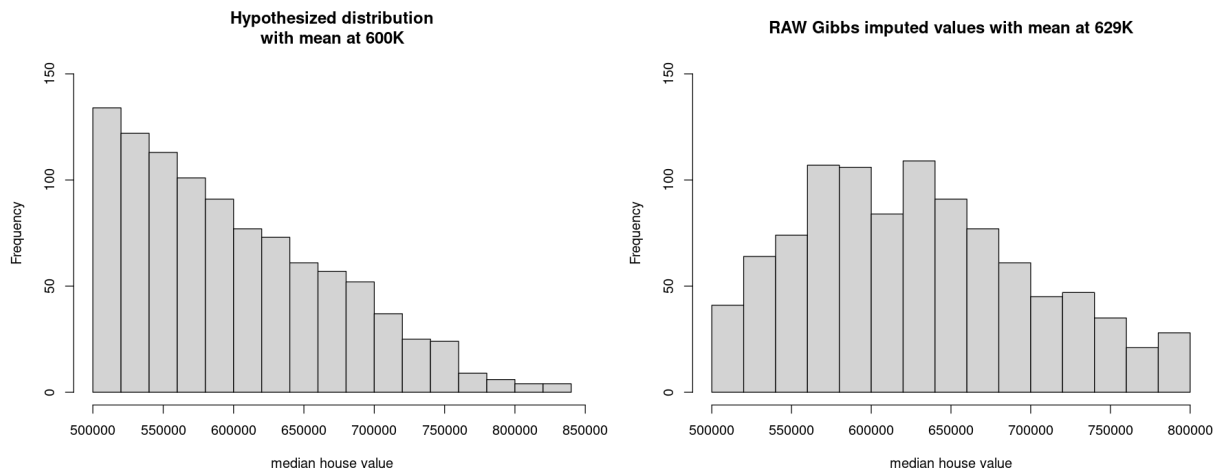
```
mat <- t(as.matrix(c(1,2)))
layout(mat, widths = rep.int(20, ncol(mat)),
       heights = rep.int(7, nrow(mat)), respect = FALSE)
```

```
# Left panel.
```

```
hist(imputed_vals_tmp, breaks=20, main="Hypothesized distribution
with mean at 600K", ylim=c(0, 150), xlab="median house value")
```

```
# Right panel.
```

```
hist(z_preds, breaks=20, main="RAW Gibbs imputed values with mean at 629K",
     ylim=c(0, 150), xlab="median house value")
```

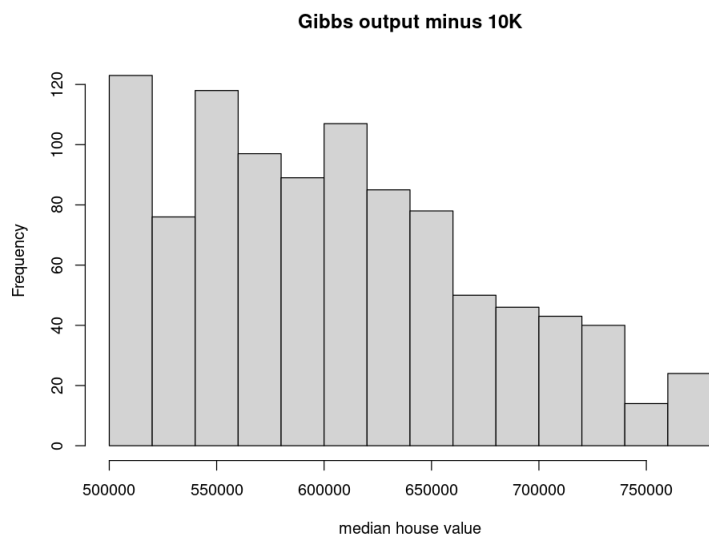


```
In [216]: # Adjust the Gibbs output so that the mean is closer to
# 605K.
```

```
z_preds_adj <- z_preds - 25000
preds_adj <- ifelse(z_preds_adj < 500000, 500000, z_preds_adj)
```


```
options(repr.plot.width= 8, repr.plot.height= 6)
```

```
hist(preds_adj, breaks=14, main="Gibbs output minus 10K",
     xlab="median house value")
```



```
In [217]: options(repr.plot.width= 8, repr.plot.height= 6)

hist(preds_adj + 100, breaks=40,
     main="Distribution of imputed median_house_values
after 20K adjustment (zoom)", xlim= c(500000, 550000),
     xlab="median house value")
```



```
In [218]: summary(preds_adj)
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
500000   548851   599113   604761   652362   774914
```

```
In [ ]: ### COMMENTS:

# The mean is now about where we have predicted it to be.

# But we do not expect there to be a sudden drop in the
# number of districts from 500K to 505K; we expect the
# drop, if there is one, to be more gradual.

# We can fix this by adjusting z_means prior to
# calling rnorm.trunc03.
```

```
In [221]: # Instead of using 605K in what follows, we will need to over-
# correct a bit.

(z_means_bar <- mean(z_means))

z_means_adj <- z_means + (590000^response_var_power - z_means_bar)
summary(z_means_adj)
round(mean(z_means_adj)^inv_pwr)
```

```
11.0306195607945
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
10.9     10.9     10.9     10.9    11.0    11.3
```

```
590000
```

```
In [225]: C_upper
```

```
11.5493142302066
```

```
In [229]: # Get some new, adjusted predictions.
```



```
# Also adjust C_upper a bit.

C_upper_adj <- 11.46

set.seed(1931)
preds_adj02 <- round(rnorm.trunc03(n.censored, z_means_adj, z_sds, lo=C, hi=C_upper_adj), 5)
preds_adj02 <- round(preds_adj02^inv_pwr)
summary(preds_adj02)

  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
500431  554334  598596  605688  648739  766141
```

```
In [ ]: ### COMMENT:
```

```
# The mean is now around where we expect it to be.
```


```
In [231]: # Compare the adjusted values to the hypothesized distribution.
```

```
options(repr.plot.width= 15, repr.plot.height= 6)

mat <- t(as.matrix(c(1,2)))
layout(mat, widths = rep.int(20, ncol(mat)),
       heights = rep.int(7, nrow(mat)), respect = FALSE)

# Left panel.
hist(imputed_vals_tmp, breaks=20, main="Hypothesized distribution
with mean at 600K", ylim=c(0, 150), xlab="median house value")

# Right panel.
hist(preds_adj02, breaks=14, main="Adjusted Gibbs imputed values with mean at 606K",
     ylim=c(0, 150), xlab="median house value")
```



```
In [ ]: ### COMMENT:
```

```
# The shape of the adjusted Gibbs output is still fairly
# different from what we hypothesize; one would expect
# more values to be found between 500K and 550K. This
# is not something we can do much about without disturbing
# the random nature of the Gibbs output. My experience is
# that if we try to manually change the shape of the
# distribution (this can be done without changing the order
# of the predictions), we adversely affect the RSS gain
# (see Appendix C for the RSS gain measure).
```

```
In [232]: # Assign imputed values.
```

```
newdat <- dat
newdat$median_house_value[censored] <- preds_adj02
```

```
summary(newdat$median_house_value)
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 15000  119600  179800  211987  264950  766141
```

In [233]: *# Plot both before and after.*


```
options(repr.plot.width= 15, repr.plot.height= 7)

mat <- t(as.matrix(c(1,2)))
layout(mat, widths = rep.int(20, ncol(mat)),
       heights = rep.int(7, nrow(mat)), respect = FALSE)

# layout.show(n = 2)

# plot the "before" scatter
plot(dat$median_income, dat$median_house_value, type= "p", pch=1, cex=0.5, col="grey",
     xlab= "median_income", ylab= "median_house_value", ylim= c(0, 0.80e06), xlim= c(0, 15),
     main= "Data prior to imputation ", cex.main=1.6)

# plot the newly predicted values
plot(newdat$median_income, newdat$median_house_value, type= "p", pch=1, cex=0.5, col="grey",
     xlab= "median_income", ylab= "median_house_value", ylim= c(0, 0.80e06), xlim= c(0, 15),
     main= "Data with adjusted imputed median house values", cex.main=1.6)
```



## Save to disk

In [234]: *# Save imputed values for median\_house\_value.*

```
write.csv(newdat,
        file="/home/greg/Documents/stat/Geron_ML/datasets/housing/housing_cleaned_v03pt5.csv",
        row.names=TRUE)
```

In [235]: `dat <- newdat`  
`rm(newdat)`

## Final Comments for Appendix B

A check to see that the imputed values are consistent with the data and with the Gibbs sampler output is made in Section 2 of Part01. Material in Appendix C also helps to make the case that this method of adjusting the Gibbs output "works" and in fact is an improvement upon the raw Gibbs output.

The Gibbs sampler method of imputation does not necessarily give us an appropriate distribution shape for the imputed values, nor even an appropriate mean. In appendices A and B I have shown how we can predict with some confidence where the mean of the imputed values ought to lie. As we will see in Appendix C, adjusting the Gibbs output to re-situate the mean can be done without adversely affecting the quality of the predictions (in terms of the RSS gain measure discussed in Appendix C).

When imputing values at the tail of a variable's distribution, the Gibbs sampler does not know that we are near the tail. This is one reason why the output from the Gibbs sampler is likely to have a mean that is too high, and it is a reason why the distribution of the imputed values is likely to not be consistent with what we expect. To correct for a mean that is too high, we can shift the Gibbs output to the left. This increases the record counts near the cap for the censored data. This is good because it causes the distribution of the Gibbs output to look more like what we expect at the tail. Beyond this, we cannot do much about the shape of the output without adversely affecting the quality of the predictions. An example of this is shown in Appendix C.

In [ ]: